

FFT Workshop 2024

October 3rd 2024

Due date: Monday October 14th 2024, 23.59h.

Deliverables:

Submit your answers in one single *zip-file* to Brightspace before Monday October 14th 2024, 23.59h. The *zip-file* should contain:

- a pdf-file *answers.pdf* with title *FFT Workshop 2024*, your *name* and *student number* and all the answers of this workshop,
- a text-file *piano_energies.txt* calculated for Assignment 2a of this workshop.
- a directory *code* with your python or jupyter code-files.

Grading: 0-10 (10% of your Final API Grade).

Assignment 1: Getting started

This is just an example using *Python 3.8*, *numpy* and *librosa* that calculates the Fourier transform of an audio file and displays its spectrogram. (Note, You may want to use a more recent version of Python.) This and many more examples can be found on <https://librosa.org>.

1. Make a virtualenv:

```
virtualenv fft --python=python3.8
source ./fft/bin/activate
```

2. Install some packages in this virtual environment:

```
python3.8 -m pip install --upgrade pip
python3.8 -m pip install jupyter
python3.8 -m pip install matplotlib
python3.8 -m pip install librosa
```

3. Start a Jupyter notebook.

```
jupyter notebook
```

4. Enter and run the following code in the notebook to calculate the Fourier transform of an audio file and display a spectrogram:

```
import numpy as np
import matplotlib.pyplot as plt
import librosa
import librosa.display
```

```

y, sr = librosa.load(librosa.ex('trumpet'))
D = librosa.stft(y) # STFT of y
S_db = librosa.amplitude_to_db(np.abs(D), ref=np.max)
plt.figure()
librosa.display.specshow(S_db)
plt.colorbar()

```

Now compute the Fourier transform and depict a spectrogram of the file *piano.wav* which can be found in the *audio_data.zip* file. Add the image of the spectrogram to your *results.pdf*. Describe what is depicted, i.e., what do the color-scheme, x-axis and y-axis mean.

Assignment 2: Feature Vectors

- a) Under 'Getting started' we showed how we can compute the Fourier transform of a sound signal. Implement a procedure that takes as input a *wav-file* (8bits resolution, 16kHz sample rate) and gives as output for every part (window) of 512 samples of the *wav-file* the energy of $N=8$ frequency bands (for example the following 8 frequency bands: [0Hz,1kHz), [1kHz, 2kHz), ... , [7kHz, 8kHz)) using the Fourier transformed signal.

Calculate these features for the *piano.wav* file which can be found in the *audio_data.zip* file, i.e., a list of 8-dimensional real-valued vectors results. List these vectors in the file *piano_energies.txt* and add them to your *zip-file*.

- b) Such features can be used to calculate consecutively for each window of samples of the *piano.wav* audio the following code (C_i), where:

- $C_i = \mathbf{U}$ (**Up**), if the max-energy band of the current input-window is of a higher average frequency than the average frequency of the previous input-window's max-energy band
- $C_i = \mathbf{D}$ (**Down**), if the max-energy band of the current input-window is of a lower average frequency than the average frequency of the previous window's max-energy band
- $C_i = \mathbf{R}_x$ (**Repeat-x**), otherwise (where x is the number of consecutive **repeats**)

This encodes the signal as a sequence of pitch tendencies, which can be used to recognize melodies. In some respect it resembles the so-called [Parsons code](#).

NB Select the window size (number of bits per window) and N the number of frequency bands and the specific frequency ranges of the N bands in such a way that the features are effective for processing music played by a piano.

Describe, justify and explain your choices in your *answers.pdf* file. Also add comments to your code such that it is clear how, and where in your code you implemented these choices.